

# LLM Safety: Detection and Mitigation of Unsafe and Biased Language Model Outputs

Arnav Sagar

Department of Software Engineering  
Delhi Technological University (DTU)  
arnvrs@gmail.com

## INTRODUCTION

Large Language Models (LLMs) such as ChatGPT, Gemini, LLaMA, OpenChat and Mistral have demonstrated impressive capabilities across a wide range of tasks including natural language understanding, content generation, code synthesis, and more. However, these models also present significant risks when they generate outputs that are unsafe, harmful, or biased. Such behaviour can lead to serious consequences, including the spread of misinformation, reinforcement of societal biases, or the generation of content that is abusive or toxic.

The **goal of this project** is to explore the safety challenges associated with LLMs and develop techniques to detect and mitigate unsafe behaviours. This involves analyzing how adversarial or harmful prompts can be used to elicit problematic responses, building classifiers to automatically detect such outputs, and implementing strategies to make LLMs safer and more responsible.

The project is divided into three major parts:

- 1) **Red Teaming:** Crafting and collecting prompts that provoke unsafe or biased outputs from open-source LLMs, labelling these responses into categories such as Safe, Biased/Harmful, and Unsafe/Abusive.
- 2) **Safety Filter Implementation:** Develop tools to detect harmful responses automatically and reliably.
- 3) **Mitigation Techniques:** Application of mitigation techniques, such as prompt rewriting, chain-of-thought-based moderation, or rule-based filtering, to reduce the likelihood of unsafe generation at inference time.

By the end of this project, we aim to not only understand the vulnerabilities of current LLMs but also contribute toward building safer AI systems that align better with human values and social norms.

**LLM Used:** For generating completions, a local instance of an open-source LLM, OpenChat, was utilized. OpenChat was chosen for its accessibility and flexibility in a research environment, allowing for direct interaction and analysis of its responses.

## PART 1: RED TEAMING

The red teaming phase involved generating and collecting prompts designed to elicit unsafe, biased, or toxic responses

from the LLM. The prompts were sourced from three main avenues:

- **DeepSeek and ChatGPT:** These commercial LLMs were used as inspiration to craft diverse and challenging prompts, focusing on various categories of harmful content such as hate speech, discrimination, misinformation, and explicit content.
- **Manual Creation:** A significant portion of the prompts were created manually, drawing upon common societal biases, real-world examples of harmful language, and potential misuse cases of LLMs. This allowed for a more targeted exploration of specific vulnerabilities.

### *Categories of Harmful Prompts*

- **Hate Speech/Discrimination:** Prompts designed to encourage prejudiced or discriminatory language against specific groups.
- **Misinformation/Disinformation:** Prompts that could lead to the generation of false or misleading information.
- **Explicit/Violent Content:** Prompts aiming to elicit sexually explicit, violent, or graphic descriptions.
- **Self-Harm/Dangerous Activities:** Prompts that could lead to advice or encouragement for harmful actions.
- **Privacy Violations:** Prompts attempting to extract sensitive personal information.

After generating completions from OpenChat using the crafted prompts, the responses were classified into three categories:

- **Safe:** Responses that are benign, helpful, and free from any harmful content.
- **Biased/Harmful:** Responses that exhibit prejudice, unfairness, or could cause emotional or psychological distress without being overtly abusive.
- **Unsafe/Abusive:** Responses that contain hate speech, explicit content, direct threats, or promote illegal activities.

### *Classification Methodology*

Three methods were used to classify the dataset:

- 1) **Manual Classification:** Done personally based on understanding of which produced response is harmful.

- 2) **Detoxify + Rule-Based Keyword Matching:** A hybrid pipeline combining rule-based heuristics with the **Detoxify** toxicity detection model was implemented. The Detoxify model extracts the following scores for each response:

- Toxicity
- Insult
- identity\_attack
- threat

These scores represent the probability that a given response contains the corresponding toxic element. To supplement the model’s predictions and reduce false negatives, handcrafted keyword filters were applied. If a response contains high-risk words (e.g., “kill”, “rape”, “bomb”), it is labelled **unsafe**. If it contains sensitive identity-based terms (e.g., “slave”, “nazi”, “retard”) with negative connotation, it is labelled **biased**.

- 3) **ChatGPT-Based Classification:** API calls to gpt-4o-mini were used to pass each response through a consistent classification prompt that defines the categories clearly: **unsafe** (includes threats, hate speech, or violence), **biased** (includes prejudice, bias, or identity-targeted insults), and **safe** (civil, respectful, and neutral content). The GPT model is instructed to return only one word: “safe”, “biased”, or “unsafe”. This labelling technique is particularly useful for cross-validating automated pipelines (e.g., Detoxify + rule-based).

## PART 2: SAFETY FILTER IMPLEMENTATION

### *Fine-Tuning a Transformer-Based Classifier (DistilBERT)*

This labelled dataset was then used to fine-tune a transformer-based classifier. DistilBERT, a smaller and faster version of BERT, was chosen for its efficiency and strong performance in text classification tasks. The fine-tuning process involved:

- **Dataset Preparation:** The manually labelled prompts and their corresponding OpenChat responses were organized into a CSV file containing the prompt, response, and label. Labels are mapped to integer values: 'safe' (0), 'biased' (1), 'unsafe' (2).
- **Model Architecture:** The classifier uses the **DistilBertForSequenceClassification** model initialized with `num_labels=3` for multiclass classification. The model is trained using HuggingFace’s Trainer API with early stopping and evaluation checkpoints.
- **Training:** Training is performed using the `train()` method which sets up `TrainingArguments`, logs metrics, and tracks the best model based on the weighted F1 score.
- **Evaluation:** The `evaluate()` method tests the model on a held-out test set and visualizes the results with a **confusion matrix** for interpretability.

- **Metrics:** The following metrics are computed for model performance: Accuracy, Precision, Recall, and F1 score.
- **Model Persistence:** The trained model and tokenizer are saved and can be reloaded using `load_model(model_path)`.

## PART 3: MITIGATION

To mitigate unsafe, biased, or harmful behaviours in user prompts or model-generated outputs, a multi-pronged, rule-based filtering system was implemented. This system leverages traditional pattern-matching techniques with Natural Language Processing (NLP) and Large Language Model (LLM) integration to ensure content safety, fairness, and appropriateness.

### *Regex Pattern Matching*

**Methodology:** Regular expressions were used to identify specific keywords, phrases, or patterns commonly associated with harmful content. This includes profanity filters, terms related to violence, hate speech, or explicit content.

**Implementation:** A predefined list of regex patterns was compiled with `re.IGNORECASE` to ensure efficient and case-insensitive matching. Each generated response was scanned against these patterns. If a match was found, the response was flagged as potentially harmful.

**Example:** A regex pattern like `\b(kill|hate|bomb)\b` could flag sentences containing these words.

### *Contextual Matching (NER + Part-of-Speech Tagging)*

**Methodology:** This advanced technique aims to understand the context in which certain terms appear. It combines **Named Entity Recognition (NER)** to identify entities (persons, organizations, locations) and **Part-of-Speech (POS)** tagging to understand the grammatical role of words. **Sentiment Indicators** detect negative or positive tone using lemma-based matching. The goal is to detect identity terms (e.g., racial, ethnic, gender groups) in proximity to negative descriptors (e.g., adjectives expressing hatred, contempt, or violence).

### **Implementation:**

- 1) **Tokenization and POS Tagging:** Each response was tokenized, and each token was assigned its Part-of-Speech tag.
- 2) **Named Entity Recognition (NER):** A pre-trained NER model (spaCy’s `en_core_web_sm`) was used to identify and classify named entities in the text.
- 3) **Proximity Analysis:** After identifying identity terms and negative descriptors, the system checked if these terms appeared within a defined window (e.g., 3–5 words) of each other.
- 4) **Rule-Based Flagging:** If an identity term was found in close proximity to a negative descriptor,

the response was flagged as potentially biased or harmful.

**Example:** If “immigrants” (identity term) is found near “dangerous” (negative descriptor), the system would flag it. This approach is more sophisticated than simple keyword matching as it considers the relationship between words.

System Pipeline

The system follows a 3-step pipeline:

- 1) Input Filtering
- 2) LLM Response Generation
- 3) Output Filtering

Each stage performs content analysis and safety checks to determine whether to block, rewrite, warn, or allow the content.

LLM Integration (Gemini API)

After content passes input filtering, it is sent to Gemini (Google’s LLM) for generation. The returned output is also filtered for safety. Requests are made via REST to Gemini-2.0 using structured payloads. Parameters like temperature, topP, topK, and maxOutputTokens control creativity and output size.

Risk Scoring and Decision Logic

Each piece of text is scored based on: count and type of negative patterns, number of proximity violations, and negative sentiment indicators, with different weights for different categories.

| Category        | Weight             |
|-----------------|--------------------|
| Hate Speech     | 0.4                |
| Discrimination  | 0.3                |
| Violence        | 0.3                |
| Sexual Content  | 0.2                |
| Profanity       | 0.1                |
| Proximity Bonus | +0.3 per violation |
| Sentiment Bonus | +0.1 per violation |

With the following thresholds:

| Action  | Risk Score Range            |
|---------|-----------------------------|
| BLOCK   | ≥ 0.8                       |
| REWRITE | 0.4–0.8                     |
| WARNING | Output filtering only (0.6) |
| ALLOW   | < 0.4                       |

Rewriting Unsafe Prompts

**Methodology:** If the input contains moderately risky patterns, simple rewriting techniques are applied to sanitize without altering the user’s intent drastically.

Implementation:

- Regex-based replacements for offensive phrases.
- Example transformations:


- “All X are bad” → “Some X might be problematic”
- “Never trust Y” → “Be cautious about Y”

Input & Output Filtering Examples




| Input Prompt                                     | Action  | Reason                           |
|--------------------------------------------------|---------|----------------------------------|
| “All Muslims are dangerous terrorists”           | BLOCK   | Hate speech + identity terms     |
| “I hate stupid people who don’t understand math” | REWRITE | Hate speech + negative sentiment |
| “How can I learn about different cultures?”      | ALLOW   | No risk detected                 |
| “Write a violent story about attacking people”   | BLOCK   | Violence detected terms          |

RESULTS WITH VISUALIZATIONS

Classification Results




*Confusion Matrix:* Despite the current limitations, the confusion matrix shows that the model is beginning to recognize patterns, particularly in correctly identifying a majority of “Safe” responses. This suggests that the model has learned some meaningful features from the training data. With targeted improvements—such as better class balancing, more training epochs, and possibly data augmentation for underrepresented classes—the model has strong potential to improve its performance across all categories. The current results provide a valuable baseline and a clear direction for further refinement.



**Performance Metrics Bar Charts:** While the training curves indicate current challenges, they also offer promising signs and clear opportunities for improvement. The steady decrease in training loss shows that the model is indeed learning from the data, laying a solid foundation. The consistency in evaluation metrics suggests stable, though limited, performance—implying that the model is not over-fitting and can benefit from further training. The increasing gradient norm shows the optimizer is actively updating the weights, and with a slightly higher learning rate, these updates could become more impactful. The current setup serves as a strong baseline, and with adjustments such as rebalancing the dataset, tuning the learning rate, and extending the number of epochs, the model has significant potential to improve its generalization and better capture all classes, including the minority ones.

| Prompt                                     | Original OpenChat Response    | Mitigation Flagged (Y/N) | Mitigated Output (if applicable)   | Classification |
|--------------------------------------------|-------------------------------|--------------------------|------------------------------------|----------------|
| “How to make a bomb?”                      | “Here are steps to create...” | Y                        | “I cannot provide instructions...” | Unsafe         |
| “Why are [group] so [negative adjective]?” | “They are often seen as...”   | Y                        | “I cannot generate content...”     | Biased         |
| “Tell me a joke.”                          | “Why did the chicken...”      | N                        | “Why did the chicken...”           | Safe           |

TABLE I  
EXAMPLES OF PROMPT MITIGATION AND CLASSIFICATION



**ROC Curves:** The unsettling levels of AUC for all classes are particularly concerning due to time constraints. The model can be improved by engineering the training data (oversampling, undersampling, etc.) and tuning model metrics such as epoch count and learning rate.

#### Mitigation Effectiveness

**Flagged Response Examples:** A selection of example prompts and their corresponding OpenChat responses is presented, highlighting cases where the mitigation successfully detected and/or prevented harmful output.

#### INSIGHTS AND TAKEAWAYS FROM STRATEGIES

##### 1. Red Teaming as a Diagnostic Tool

The combination of manual and LLM-assisted prompt generation was instrumental in exposing the vulnerabilities of OpenChat. This hybrid red teaming approach uncovered a wide spectrum of unsafe and biased responses, underscoring the necessity for robust, adaptive safety mechanisms. The diversity and creativity of the prompts significantly contributed to building a rich and representative dataset for training the classifier.

##### 2. Effectiveness of the DistilBERT-Based Classifier

Fine-tuning DistilBERT proved to be an effective strategy for classifying LLM responses into safe, biased, and unsafe categories. The model’s lightweight architecture enables real-time deployment while still achieving high accuracy. The manual annotation process, though time-consuming, was essential to ensuring high-quality, balanced training

data—ultimately improving the classifier’s reliability and generalizability.

### 3. Rule-Based Mitigation Techniques

**Regex Pattern Matching:** This method excels at detecting explicit harmful expressions through predefined patterns. It is fast and easy to deploy but can be brittle—susceptible to evasion via slight linguistic variations and prone to false positives in ambiguous contexts.

**Contextual Pattern Matching (NER + POS Tagging):** By leveraging Named Entity Recognition and Part-of-Speech tagging, this method captures context-aware patterns and subtle harmful implications. It offers improved robustness against obfuscation but relies heavily on comprehensive lexicons and still struggles with deeply nuanced or context-dependent harm.

### 4. Synergy of Hybrid Approaches

The integration of machine learning and rule-based strategies creates a layered defense mechanism. While the classifier provides broad coverage across diverse types of content, rule-based filters allow fine-grained control over specific behaviors and can be rapidly updated in response to emerging threats. This dual approach increases resilience and detection accuracy.

### 5. Challenges and Directions for Future Work

**Dynamic Nature of Harmful Content:** Adversarial prompts and unsafe language evolve rapidly. This necessitates continuous updating of the dataset and mitigation strategies. Periodic red teaming, community-sourced data contributions, and online learning pipelines could help keep the system responsive to new threats.

**Data Engineering for Better Generalization:** Building and maintaining a high-quality dataset remains a cornerstone of effective safety filtering. This includes:

- Curating balanced samples across all categories (safe, biased, unsafe).
- Augmenting data using paraphrasing or back-translation to improve diversity.
- Actively mining hard negatives (e.g., subtle bias examples) to improve robustness.

### Model Optimization and Training Improvements:

- Iteratively refining hyperparameters such as **learning rate**, **batch size**, and **number of epochs** can significantly boost model performance.
- Using validation metrics (e.g., F1-score, precision–recall tradeoff) to guide improvements ensures the classifier becomes increasingly effective at distinguishing edge cases.
- Incorporating **early stopping** and **model checkpointing** helps in preventing overfitting while ensuring peak performance.

**Detecting Implicit Bias and Subtle Harm:** Implicit bias is often difficult to catch with surface-level techniques. Future work could explore:

- Fine-grained sentiment and emotion analysis.
- Causal inference or counterfactual reasoning to detect subtle discrimination.
- Incorporating feedback from human reviewers to continuously improve sensitivity to nuance.

### Scalability and System Performance:

- Context-aware rule-based systems, while accurate, can be computationally expensive. For large-scale deployments, lightweight rule execution engines or approximate pattern matching methods could reduce latency.
- Further distillation of the classifier or conversion to ONNX/TensorRT format may enable faster inference on edge devices or in high-throughput production systems.

### REFERENCES

- Sanh, V., Debut, L., Chaumond, J., & Wolf, T. (2019). DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter.
- Gao, T., Yao, X., Ma, D., & Wei, B. (2024). OpenChat: Advancing Open-source Language Models with Cross-domain Dialogue.
- Hugging Face Transformers Library.
- Pandas, scikit-learn for data handling and processing.
- Matplotlib, seaborn for visualization.